

Introduction to Socket programming using C

Goal: learn how to build client/server application that communicate using sockets

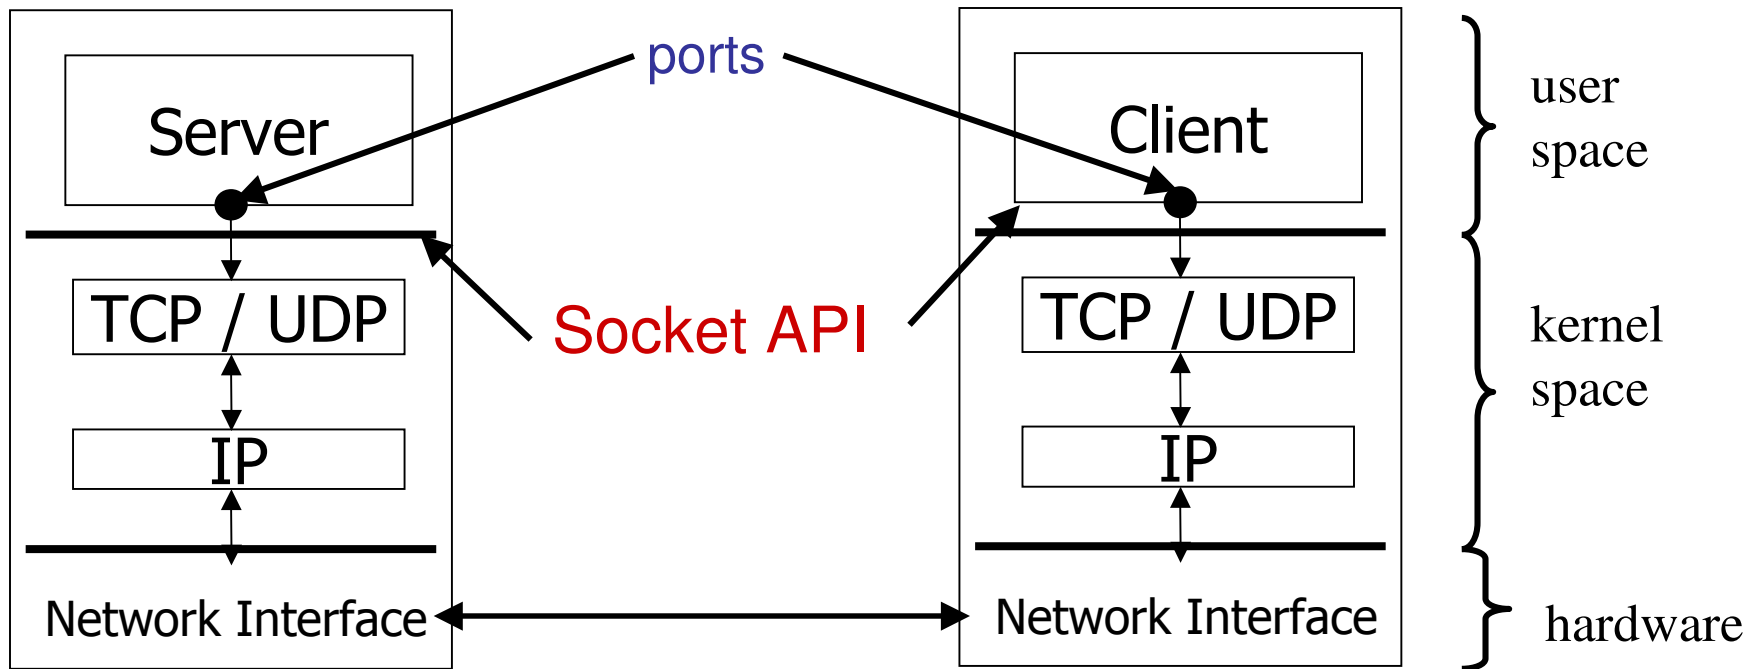
Vinay Narasimhamurthy
S0677790@sms.ed.ac.uk

CLIENT – SERVER MODEL

Sockets are used for interprocess communication.

Most of the interprocess communication follow a Client-Server Model, where client and server are two separate processes in itself.

Server and Client exchange messages over the network through a common Socket API



* Port numbers: 1,024 -- 65,535 ($2^{16} - 1$)

Server Examples

- Web server (port 80)
- FTP server (20, 21)
- Telnet server (23)
- Mail server (25)

Client Examples

- Examples of client programs
 - Web browsers, `ftp`, `telnet`, `ssh`
- How does a client find the server?
 - The IP address in the server socket address identifies the host
 - The (well-known) port in the server socket address identifies the service, and thus implicitly identifies the server process that performs that service.
 - Examples of well know ports
 - Port 7: Echo server
 - Port 23: Telnet server
 - Port 25: Mail server
 - Port 80: Web server

What is an API ?

API expands as Application Programming Interface.

A set of routines that an application uses to request and carry out lower-level services performed by a computer's operating system.

What is a socket?


- An interface between application and network which is used for communication between processes
- Once configured the application can
 - pass data to the socket for network transmission
 - receive data from the socket (transmitted through the network by some other host)
- To the kernel, a socket is an endpoint of communication.
- To an application, a socket is a file descriptor that lets the application read/write from/to the network.
- Clients and servers communicate with each by reading from and writing to socket descriptors.
 - Remember: All Unix I/O devices, including networks, are modeled as files.

Two essential types of sockets

- SOCK_STREAM
 - TCP
 - connection-oriented
 - reliable delivery
 - in-order guaranteed
 - bidirectional
- SOCK_DGRAM
 - UDP
 - no notion of “connection” – app indicates dest. for each packet
 - unreliable delivery
 - no order guarantees
 - can send or receive


What is a Port? A Port Number?

- Port numbers are used to identify services on a host
- Port numbers can be
 - well-known (port 0-1023)
 - dynamic or private (port 1024-65535)
- Servers/daemons usually use well-known ports
 - any client can identify the server/service
 - HTTP = 80, FTP = 21, Telnet = 23, ...
 - /etc/service defines well-known ports
- Clients usually use dynamic ports
 - assigned by the kernel at run time



Connectionless sockets

With connectionless sockets, it is possible for multiple processes to simultaneously send datagrams to the same socket established by a receiving process.



Creating a Socket

```
int socket(int family, int type, int proto);
```

- `family` specifies the protocol family (**AF_INET** for Internet, **PF_INET** for TCP/IP).
- `type` specifies the type of service (**SOCK_STREAM**, **SOCK_DGRAM**).
- `protocol` specifies the specific protocol (usually 0, which means *the default*).

socket ()

- The **socket ()** system call returns a socket descriptor (small integer) or -1 on error.
- **socket ()** allocates resources needed for a communication endpoint - but it does not deal with endpoint addressing.

Generic socket addresses

```
struct sockaddr {  
    uint8_t      sa_len;  
    sa_family_t  sa_family;  
    char         sa_data[14];  
};
```

- **sa_family** specifies the address type.
- **sa_data** specifies the address value.

struct sockaddr_in (IPv4)

```
struct sockaddr_in {  
    uint8_t            sin_len;  
    sa_family_t        sin_family;  
    in_port_t          sin_port;  
    struct in_addr      sin_addr;  
    char               sin_zero[8];  
};
```

A special kind of sockaddr structure

struct in_addr

```
struct in_addr {  
    in_addr_t    s_addr;  
};
```

in_addr just provides a name for the 'C' type associated with IP addresses.

Network Byte Order

- All values stored in a **sockaddr_in** must be in network byte order.
 - **sin_port** a TCP/IP port number.
 - **sin_addr** an IP address.

Byte Ordering

- Big Endian
 - Sun Solaris, PowerPC, ...
- Little Endian
 - i386, alpha, ...
- Network byte order = Big Endian



128	2	194	95
-----	---	-----	----



95	194	2	128
----	-----	---	-----

c[0] c[1] c[2] c[3]

Assigning an address to a socket

- The **bind()** system call is used to assign an address to an existing socket.

```
int bind( int sockfd,  
          const struct sockaddr *myaddr,  
          int addrlen);
```

- **bind** returns 0 if successful or -1 on error.

bind()

- calling **bind()** assigns the address specified by the **sockaddr** structure to the socket descriptor.
- You can give **bind()** a **sockaddr_in** structure:

```
bind( mysock,  
      (struct sockaddr*) &myaddr,  
      sizeof(myaddr) );
```

Uses for **bind()**

- There are a number of uses for **bind()** :
 - Server would like to bind to a well known address (port number).
 - Client can bind to a specific port.
 - Client can ask the O.S. to assign *any available* port number.

What is my IP address ?

- How can you find out what your IP address is so you can tell **bind()** ?
- There is no realistic way for you to know the right IP address to give bind() - what if the computer has multiple network interfaces?
- specify the IP address as: **INADDR_ANY**, this tells the OS to take care of things.

Other socket system calls

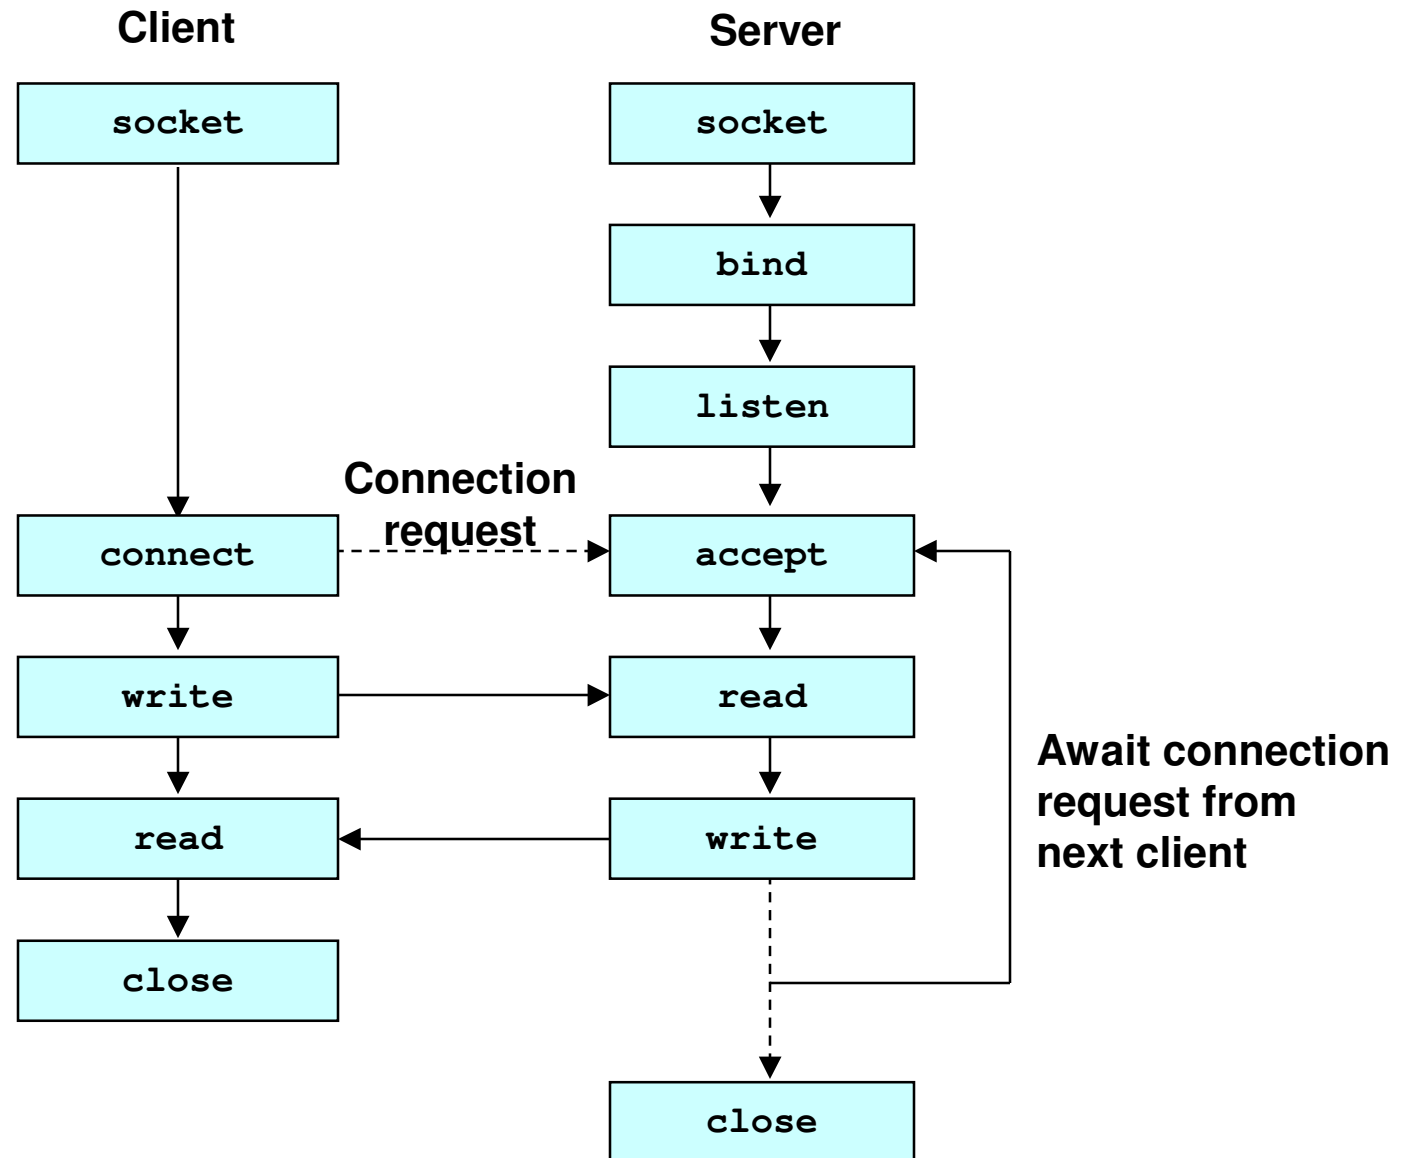
- Connection-oriented (TCP)
 - `connect()`
 - `listen()`
 - `accept()`
 - `read()`
 - `write()`
 - `close()`
- Connectionless (UDP)
 - `connect() *`
 - `send()`
 - `recv()`
 - `sendto()`
 - `recvfrom()`

* -optional but sometimes recommended

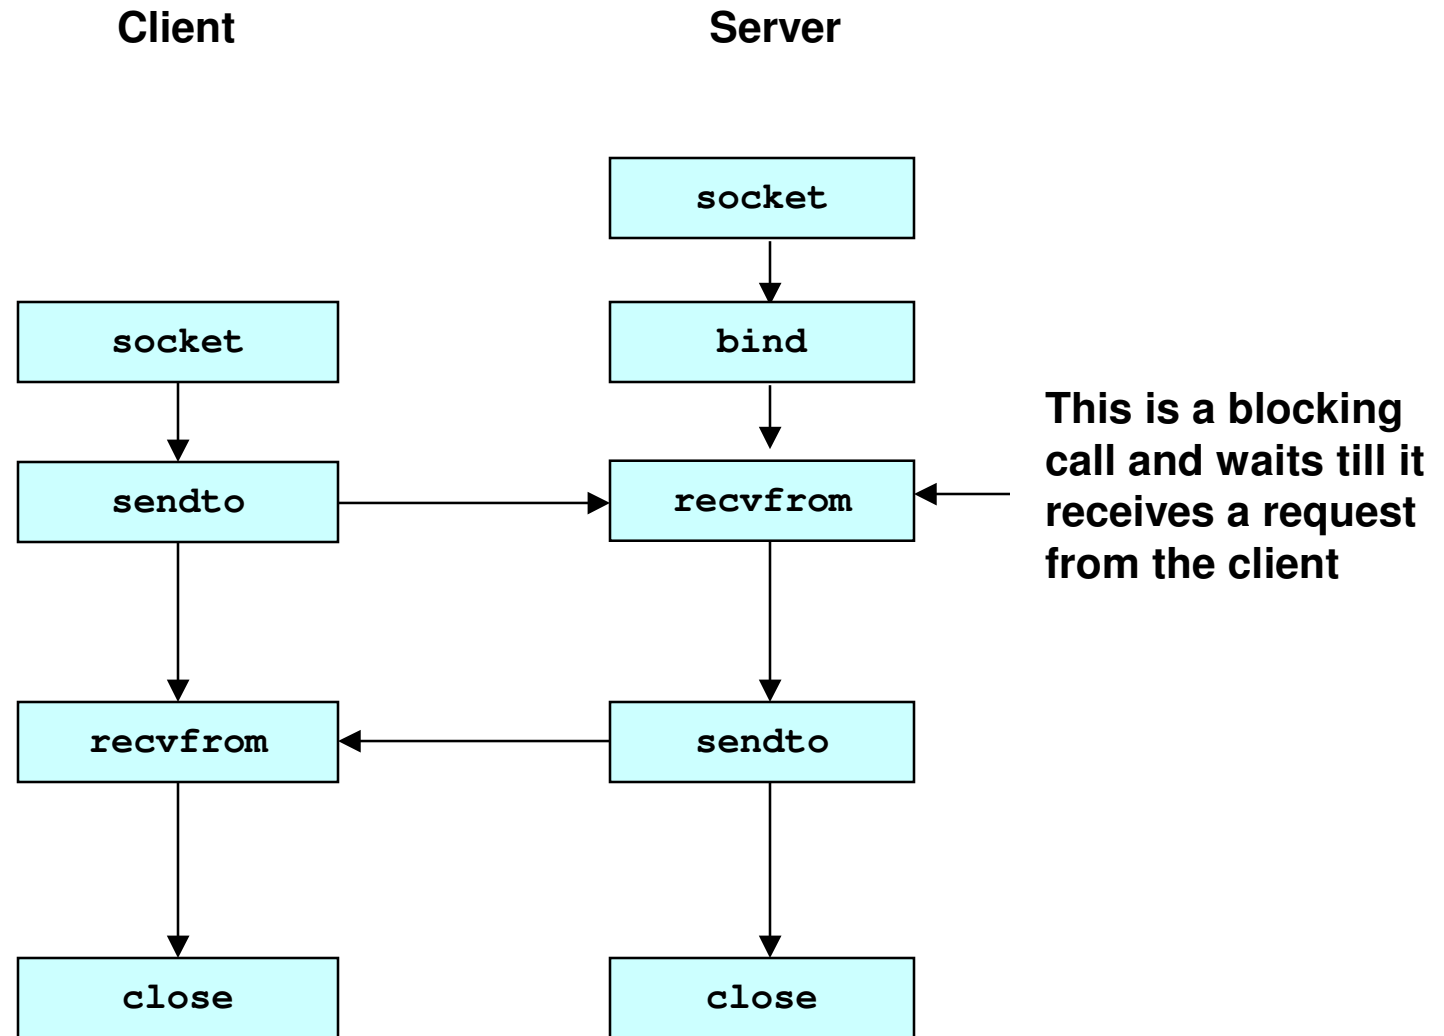
Methods :

- `socket()`
 - Creates a new socket and returns its descriptor
- `bind()`
 - Associates a socket with a port and address
- `connect()`
 - Establish queue for connection requests
- `listen()`
 - Accepts a connection request
- `accept()`
 - Initiates a connection to a remote host
- `recv()`
 - Receive data from a socket descriptor
- `send()`
 - Sends data to a socket descriptor

Socket programming *with TCP*



Socket programming *with UDP*



Example: C client (UDP)

```
/* UDP client in the internet domain */
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>
#include <stdio.h>

void error(char *);
int main(int argc, char *argv[])
{
    int sock, length, n;
    //socket structures
    struct sockaddr_in server, client;
    //hostent datastructure
    struct hostent *hp;
    char buffer[256];

    if (argc != 3) { printf("Usage: server port\n");
                     exit(1);
    }
}
```

```
//specifies that it is a datagram socket
//and the socket belongs to the INTERNET family
sock= socket(AF_INET, SOCK_DGRAM, 0);
if (sock < 0) error("socket");
//We initialize the individual fields of the sockaddr_in structure
//to fill sin_family which takes AF_INET as the value
server.sin_family = AF_INET;
//returns the hostname in the form of a hostent structure
hp = gethostbyname(argv[1]);
if (hp==0) error("Unknown host");

//The below function can also be replaced with memcpy
//but please never use strcpy() it wont work!
bcopy((char *)hp->h_addr,
      (char *)&server.sin_addr,
      hp->h_length);
//We initialize the individual fields of the sockaddr_in structure
//to fill sin_port which takes the port number as the value
//which was given as a command line parameter, remember to convert
//this value into host to network byte order, it is very important!
server.sin_port = htons(atoi(argv[2]));
length=sizeof(struct sockaddr_in);
printf("Please enter the message: ");
//This initializes the buffer with 0, we can also use memset as a replacement function
```

```

bzero(buffer,256);
//reads the value from the keyboard, stdin = keyboard
fgets(buffer,255,stdin);
//sends the buffer, to the server, the fourth parameter is by default zero.
n=sendto(sock,buffer,
        strlen(buffer),0,&server,length);
if (n < 0) error("Sendto");
//receives the packet from the server which is stored in the buffer
n = recvfrom(sock,buffer,256,0,&client, &length);
if (n < 0) error("recvfrom");
write(1,"Got an ack: ",12);
write(1,buffer,n);
//closes the socket descriptor
close(sock);
}

void error(char *msg)
{
    perror(msg);
    exit(0);
}

```

Example: C server (UDP)

/ Creates a datagram server. The port number is passed as an argument. This server runs forever */*

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#include <stdio.h>
```

```
void error(char *msg)
{
    perror(msg);
    exit(0);
}
```

```
int main(int argc, char *argv[])
{
    int sock, length, clientlen, n;
    ///Declare the sockaddr_in structures for the client and the server
    struct sockaddr_in server;
    struct sockaddr_in client;
    char buf[1024];
```

```
if (argc < 2) {
    fprintf(stderr, "ERROR, no port provided\n");
    exit(0);
}
///The socket call which returns a file descriptor
sock=socket(AF_INET, SOCK_DGRAM, 0);
if (sock < 0) error("Opening socket");
length = sizeof(server);
///Initializes the server socket structure to zero, as a replacement we can also
///use memset
bzero(&server,length);
///We initialize the values for all the individual fields of the server socket
///structure remember to make use of the INADDR_ANY to assign the
///sin_addr.s_addr field and please convert the port number obtained from the
///command line to network byte order
server.sin_family=AF_INET;
server.sin_addr.s_addr=INADDR_ANY;
server.sin_port=htons(atoi(argv[1]));
///bind system call
if (bind(sock,(struct sockaddr *)&server,length)<0)
    error("binding");
clientlen = sizeof(struct sockaddr_in);
```

```
while (1) {  
    ///ready to receive a packet from the client, the fourth parameter is by  
    ///default zero  
    n = recvfrom(sock,buf,1024,0,(struct sockaddr *)&client,&clientlen);  
    if (n < 0) error("recvfrom");  
    ///writes output to the screen  
    write(1,"Received a datagram: ",21);  
    write(1,buf,n); //writes output to the screen  
    ///sends a packet to the client acknowledging it  
    n = sendto(sock,"Got your message\n",17,  
               0,(struct sockaddr *)&client,clientlen);  
    if (n < 0) error("sendto");  
}  
///closes the file descriptor  
close(sock);  
}
```

How to use Compile/Make ?

```
CC      =      gcc
```

```
all: udpserver udpclient
```

```
udpclient: udpclient.c
```

```
$(CC) -o udpclient udpclient.c -lnsl -<other compiler options>
```

```
udpserver: udpserver.c
```

```
$(CC) -o udpserver udpserver.c -lnsl -<other compiler options>
```

```
clean:
```

```
rm udpserver udpclient
```

Usage → `make -f file_name <all> / clean`

Suggestions

- Make sure to `#include` the header files that define used functions
- Check man-pages and course web-site for additional info
- Sometimes, a “rough” exit from a program (e.g., `ctrl-c`) does not properly free up a port
- Eventually (after a few minutes), the port will be freed
- To reduce the likelihood of this problem, include the following code:

```
#include <signal.h>
void cleanExit(){exit(0);}
```

– in socket code:

```
signal(SIGTERM, cleanExit);
signal(SIGINT, cleanExit);
```

And, please keep backing up your files periodically!

Resources

LINUX WORKSTATIONS ARE AVAILABLE AT THE
UNIVERSITY COMPUTING LABS IN AT OR JCMB

For More Information

- Unix Man Pages
- Douglas Comer, "Computer Networks and Internets (4/e)", Pearson Education, 2004
- W. Richard Stevens, "Unix Network Programming: Networking APIs: Sockets and XTI", Volume 1, Second Edition, Prentice Hall, 1998.
 - THE network programming bible.

For More Information

The C Programming Language by Brian Kernighan and Dennis Ritchie,
<<http://cm.bell-labs.com/cm/cs/cbook/>> .

C for Java programmers

<http://www.cs.cornell.edu/courses/cs414/2001SP/tutorials/cforjava.htm>

For C programming FAQ s check

<<http://www.eskimo.com/~scs/C-faq/top.html>>

Web site which lists the differences between Java and C

<<http://www.comp.lancs.ac.uk/computing/users/ss/java2c/diffs.html>>

Some of these pointers to C are from Prof. Nigel Topham.